

Program Analysis for Weak Memory Models

Attila Zoltán Printz
printza@student.ethz.ch

Supervisor: Andrei Dan
Professor: Martin Vechev

Software Reliability Lab
Department of Computer Science
ETH Zurich

August 2015

Abstract

A vital use of program analysis is correctness proving. Assuming we have a program and a set of criteria our program has to meet under any circumstance, proving that our program meets all of our criteria is an important task for critical software components. Additionally, if we expand our validity requirements to concurrent runs of one or more instances of the program, which itself might have several concurrent processes, our analysis threatens to become untractably complex. Abstracting our program to an equivalent program operating on boolean variables representing our predicates, and then inputting this boolean program into a model checker, along with the requirement that all predicates must hold true, is a feasible method to prove correctness within a reasonable error margin, as shown by Ball et al. [citation]. If we consider our program to apply a weak memory model (henceforth referred to as WMM) however, this abstraction alone doesn't cover the additionally complex behaviour exhibited by WMM. To the end of being able to analyse and prove the correctness of WMM programs, we present an approach to convert a WMM program with possible non-deterministic qualities and ambiguities depending on the current memory ordering into a boolean program verifiable by our model checker. We will consider partial store order (PSO), and total store order (TSO) as possible WMM approaches.

1 Introduction

1.1 Proof of correctness

Let P be a multi-process program using a WMM instruction set, and let E be a list of predicates on the variables of P . We consider P to be correct as defined by E , iff for every possible run scenario of any number of concurrent instances of P , in which any number of processes of P may run concurrently, all predicates in E hold true in every program point of P .

By abstracting P and E as a boolean program, such a proof becomes feasible using a model checker.

1.2 Weak memory models

In all of our programs, we will differentiate between local resources, declared as local variables, and their remote counterparts. The main feature of WMM languages is the differentiation between local assignments and remote reads and writes. We consider local variables as being local both in the semantic sense (i.e. a process constitutes the scope of a local variable) and the technical sense (i.e. a local variable is stored in fast-access memory, such as a register or local cache space). In contrast, remote variables are globally shared across all processes and are assumed to be stored in shared memory, that's more time-consuming to access. We will continue using the words "local" and "remote" in this manner throughout this work.

Thus, the main advantage of WMM programming becomes clear, as the programmer will now be able to avoid costly remote memory accesses for operations restricted to one process, along with any adverse side-effects depending on the caching policy. This does not, however, mean that a remote statement will always access remote memory. For practical reasons, remote writes are immediately stored in a buffer, and only committed to remote memory in the case of a flush statement, which, as apparent from its definition, is inherently non-deterministic. In the same vein, remote reads will read from the local buffer if possible, and only access remote memory if the local buffer is empty. In this work, we will abstract the buffer as a series of local variables representing a FIFO queue. The buffering method depends on the memory model. In the case of TSO, every global variable in every process will have a dedicated buffer, whereas in the case of PSO, every process will have one buffer shared by all global variables.

1.3 Difficulties of WMM

Our program behaves differently depending on the memory ordering chosen, without this difference being readily apparent from the code. More precisely, when using TSO, the available buffer size defines a set maximum buffer size for

every global variable per process, whereas when using PSO, it only defines the sum of the buffer sizes of every global variable per process. To illustrate this with an example, if we have a program with 2 global variables x and y , with 4 remote writes in x and 2 remote writes in y before flushing, a per-process buffer size of 6 wouldn't result in a buffer overflow, as the sum of the maximum buffer sizes of x and y does not exceed 6. However, the same code results in an overflow if we define a maximum per-process per-variable buffer size of 3, since the 4th write in x would exceed the maximum buffer size allocated to x in that process.

We tackle this problem by converting our program P from one generally exploiting all the features available in WMM into another program P' that lacks any "true" WMM statements. I.e., our resulting program will only contain local assignments, remote reads, and remote writes, with no fences or flushes. Though remote reads and writes will still be featured in P' , they will simply describe immediate and deterministic remote access statements, without any of the advantages or pitfalls of their WMM counterparts. Finally, we can convert P' into a boolean program using predicate abstraction, which, as described in 1.1, can be used to prove the correctness of P regarding E .

2 Related work

This work is based on the work of Dan, Meshman, Vechev, and Yahav on predicate abstraction for relaxed memory models [1], and it implements most of the techniques derived by the same. It also borrows heavily from the predicate abstraction rules defined by Ball, Majumdar, Millstein, and Rajamani [2]. Below, we describe these publications and the nature of their influence on this work, apart from general inspiration.

2.1 Predicate Abstraction for Relaxed Memory Models

The work by Dan et al. [1] has presented the two abstraction steps this work also employs. It has handled the problem of verifying a program P running on an arbitrary memory model M by embedding a description of said memory model into a program P_M , which is in its behaviour equivalent to P . The resulting abstraction P_M has then been analysed for predicates to be generated in order to be used as correctness criteria for a successive predicate abstraction, which would then describe the conversion of P_M into a boolean program ready to be verified by a model checker. Furthermore, the aforementioned methods have been implemented for TSO and PSO.

It is safe to say that this work primarily seeks to implement the ideas and concepts described in [1], to reproduce its success in achieving feasibility in its implementation, and to improve on its performance and complexity wherever possible.

2.2 Automatic Predicate Abstraction of C Programs

The work by Ball et al. [2] has laid the foundation on which predicate abstraction in this work is implemented. It has described the difficulties model checkers face when confronted with systems with large state spaces, such as software, which is infinite-space. The authors have described the first algorithm to automatically construct a predicate abstraction of programs in C, and they have implemented it in a toolkit, which combined it with other methods such as model checking to statically verify input programs.

The final stage of this work, predicate abstraction, implements the algorithm defined by Ball et al., adjusted for the PSO/TSO language we use, and employs some of the optimizations described by the authors.

3 Programming language

We define a programming language that is both simple and powerful enough to conclusively describe weak memory models to the extent they may be available to the programmer. Throughout this work, we will refer to this language as BITCHAKU. A BITCHAKU program has only integer-type variables, but it can perform comparisons that return boolean values, which in turn can be operated on using boolean operators. Boolean values are used in conditionals and may have the values *true*, *false*, or ***, all of which are also valid literals. Comments are defined as in C, with *//* prepending line comments and */** and **/* enclosing block comments.

In this section, we will group code placeholders using *[]*-brackets, which might have a *?* appended for optional code, or *** for an arbitrary number of the code specified.

3.1 Accepting program state

A valid BITCHAKU program must have the following structure:

```
beginit
    [store statement]*
endinit

[process declaration]*

[assertion]?
```

3.2 Assertion

The assertion may be appended to the end of the program. Its syntax is derived from the boolean program syntax used by the Fender [3] model checker. The

boolean expression therein may only refer to global variables, constants, or program counters of the form `pc(i)`, with *i* being a valid process index. The value of a program counter may be that of any valid label of the process *i*.

```
assert (always (boolean expression))
```

3.3 Process declaration

Every process must have a unique integer identifier *i*.

```
process i:  
    [statement]*
```

3.4 Statement

A *statement* can be any of the following:

1. *store statement*
2. *load statement*
3. *local assignment statement*
4. *label*
5. *goto statement*
6. *flush statement*
7. *fence statement*
8. *abort statement*
9. *nop statement*
10. *if – else block*

3.5 Store statement

Generally, the store statement writes the right-hand side expression to the buffer, pending commitment to remote memory. In the special case where flushes are to be excluded, as it will be the case later on in this work, the store statement may be viewed as an immediate and deterministic remote write operation. The left-hand side identifier must be a global variable.

```
store variable identifier = integer expression;
```

3.6 Load statement

Generally, the load statement reads the most recent relevant value from the buffer, and it only reads from remote memory, if the buffer doesn't contain any values of the desired variable. In the special case where flushes are to be excluded, as it will be the case later on in this work, the load statement may be viewed as an unconditional remote read operation. The left-hand side identifier must be a local variable, while the right-hand side identifier must be a global variable.

```
load variable identifier = variable identifier;
```

3.7 Local assignment statement

The left-hand side identifier must be a local variable.

```
variable identifier = integer expression;
```

3.8 Label

The statement following the colon may not be another label. The label value i must be unique within its process.

```
 $i$ : statement
```

3.9 Goto statement

The label value i must be a valid label within the calling process.

```
goto  $i$ ;
```

3.10 Flush statement

Flushing will iterate through all global variables and non-deterministically commit the first value of the respective variable to remote memory and remove that value from the buffer. If a variable does not have any values to be committed, it will be ignored.

```
flush;
```

3.11 Fence statement

Without any operations on the buffer or memory, fences just assume the buffer to be empty from then on, thus discarding any values that might be buffered still.

```
fence;
```

3.12 Abort statement

Aborts the program with the *message* argument as the error message.

```
abort("message");
```

3.13 Nop statement

Does nothing.

```
nop;
```

3.14 If-else block

```
if (boolean expression)
    [statement]*
[else block]?
endif;
```

with the else block being:

```
else
    [statement]*
```

4 Buffer size analysis

As we have seen, some of the statements available to the programmer make use of a write buffer. Since depending on the memory model chosen, we may encounter different behaviours when running the same program using different memory models, we face the challenge of rewriting the program in a way that its behaviour will be conclusively defined, independent of the memory model chosen. We will do this by completely eliminating any buffer access by abstracting the buffers as sets of local variables. In order for us to be able to assume the lack of buffer operations, we will need to eliminate fence- and flush statements. As for remote stores and loads, we can replace them with code that while simulating their behaviours using non-WMM statements, but still including store and load statements that represent assignment statements involving global variables. Since this intermediary code will never be run, but instead will be further abstracted in the predicate abstraction stage, we may safely ignore any regular WMM-behaviour such statements might have, as we will handle store- and load statements like local assignments for predicate abstraction purposes. The


reason we refrain from using local assignment statements instead of store- or load statements in our abstracted program is our desire to preserve syntactic correctness even in our intermediary program, and to highlight the fact, that at certain points, load statements will actually perform a remote read, and that flush statements may at one point actually perform remote writes.

We define the following abstractions for TSO, and PSO respectively. Note, that for the purpose of determining the maximum buffer space needed by any global variable, we count the number of store statements in the control flow of the program within each process i . In case there are less store statements on a variable than the user-specified maximum buffer size K , we will only allocate the necessary number of buffer variables, which is the minimum of K and the number of stores on that variable in the process i . We will refer to this number as $s_{x,i}$ for any global variable \mathbf{x} .

4.1 TSO abstraction

Using TSO, every global variable has a separate buffer per process of the size $s_{x,i}$. Therefore, for each process i , for each global variable \mathbf{x} , there are $s_{x,i}$ buffer variables of the form $\mathbf{x_j_i}$, where $j \in [1, s_{x,i}]$. Since flushes are non-deterministic, we cannot keep track of occupied buffer spaces. Therefore, we declare 2 additional auxiliary variables for each global variable in each process: $\mathbf{x_cnt_i}$, which is initialised to 0 and otherwise has the value of the last buffer index to which a value of \mathbf{x} was written, and $\mathbf{x_fst_i}$, which is analogous to $\mathbf{x_cnt_i}$ with the difference that it points to the first such index. Figure 1. illustrates the buffer allocation abstraction we apply for TSO programs.

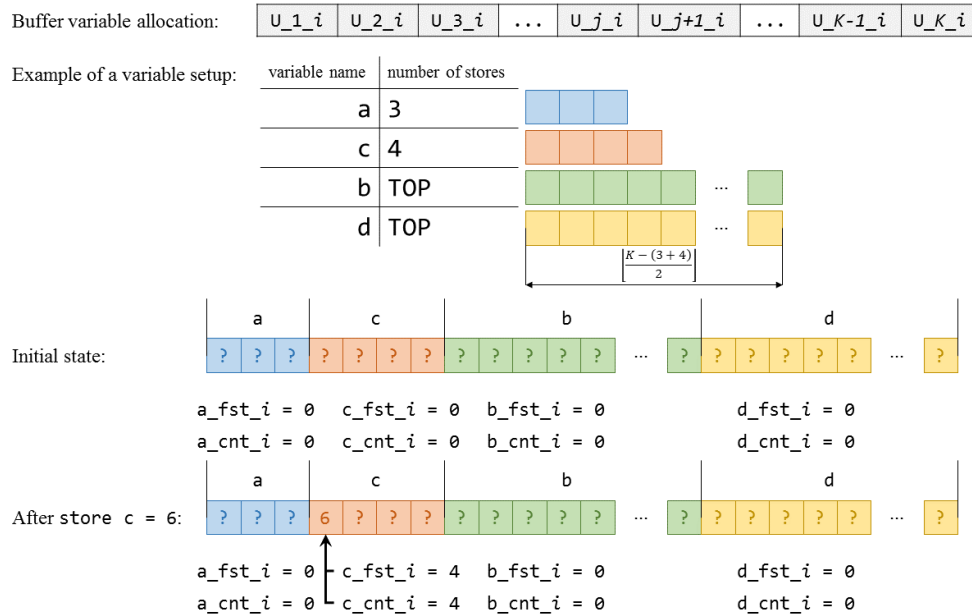
Figure 1: TSO buffer abstraction approach



4.2 PSO abstraction

Using PSO, every process has a separate buffer of the size K , which all global variables share. Therefore, it is necessary to allocate parts of the shared buffer among the global variables. After determining $s_{x,i}$ for each x , we will try to allocate the buffer space of the process i for each variable. First, we allocate buffer space for every global variable x with $s_{x,i} < K$, then we distribute the remaining variables over the remaining buffer space. This will yield us statically known offsets $o_{x,i}$ for each global variable. Thus, we can use the same approach as we did with TSO, with the added overhead of the smallest valid value of x_cnt_i and x_fst_i being $o_{x,i}$ and instead of 1, and the largest valid value of the same being 1 less than the offset of the next allocated variable, or K if x is the last one allocated. Given all $s_{x,i}$, all of these offsets are statically computable in linear time. Figure 2. illustrates the buffer allocation abstraction we apply for PSO programs.

Figure 2: PSO buffer abstraction approach



4.3 Replacement rules

After we compute all $s_{x,i}$ for each global variable x and for each process i , we have enough data to statically perform all necessary replacements in our program to allow predicate abstraction. In case we use PSO, all $o_{x,i}$ may also be computed now. We will use the notation $p_{x,i}$ to signify the lowest valid buffer index (i.e. 1 for TSO, $o_{x,i}$ for PSO), and $q_{x,i}$ to signify the largest valid buffer index (i.e. $o_{y,i}$ for the next allocated global variable for PSO, or K for TSO and the last allocated global variable for PSO).

4.3.1 Replacing store statements

Assuming \mathbf{x} is a global variable, \mathbf{r} is an integer expression, and we encounter `store $\mathbf{x} = \mathbf{r}$` in a process i , we will replace this statement with:

1. TSO approach:

```

if (x_cnt_i =  $s_{x,i}$ )
    abort("overflow");
endif;

if (x_fst_i = 0)
    x_fst_i = 1;

```

```

        x_1_i = r;
    endif;

    x_cnt_i = x_cnt_i + 1;

    if (x_cnt_i = 2)
        x_2_i = r;
    endif;
    ...
    if (x_cnt_i = j)
        x_j_i = r;
    endif;
    ...
    if (x_cnt_i = s_{x,i})
        x_{s_{x,i}}_i = r;
    endif;

```

2. PSO approach:

```

    if (x_cnt_i = q_{x,i})
        abort("overflow");
    endif;

    if (x_fst_i = 0)
        x_fst_i = p_{x,i};
        x_cnt_i = p_{x,i};
        x_{p_{x,i}}_i = r;
    else
        x_cnt_i = x_cnt_i + 1;
    endif;

    if (x_cnt_i = p_{x,i} + 1)
        x_{p_{x,i} + 1}_i = r;
    endif;
    ...
    if (x_cnt_i = j)
        x_j_i = r;
    endif;
    ...
    if (x_cnt_i = q_{x,i})
        x_{q_{x,i}}_i = r;
    endif;

```

4.3.2 Replacing load statements

Assuming x is a global variable, l is a local variable, and we encounter `load l = x` in a process i , we will replace this statement with:

```
if (x_cnt_i = 0)
    load l = x;
endif;

if (x_cnt_i = px,i)
    l = xpx,i-i;
endif;

...

if (x_cnt_i = j)
    l = xj-i;
endif;

...

if (x_cnt_i = qx,i)
    l = xqx,i-i;
endif;
```

4.3.3 Replacing fence statements

Assuming we encounter `fence`; in a process i , we will replace this statement with the following block for every global variable x :

```
assume (x_cnt_i = 0);
assume (x_fst_i = 0);
```

4.3.4 Replacing flush statements

Assuming we encounter `flush`; in a process i , we will replace this statement with:

```
p: if (*)
    [global variable flushes]
    goto p;
endif;
```

The placeholder *[global variable flushes]* describes the following code block for each global variable x :

```
if (x_cnt_i > 0)
    if (*)
        if (x_fst_i > px,i)
            if (x_fst_i > px,i + 1)
                ...
            
```

```

        if (x_fst_i > j)
            ...
            if (x_fst_i > qx,i - 1)
                store x = xqx,i-i;
            else
                store x = xqx,i-1-i;
            endif;
            ...
        else
            store x = xj-i;
        endif;
        ...
    else
        store x = xpx,i+1-i;
    endif;
else
    store x = xpx,i-i;
endif;

    x_fst_i = x_fst_i + 1;
endif;
endif;

```

5 Predicate abstraction

After the replacements carried out in the buffer size analysis section, we have a program that both encapsulates the behaviour of the chosen WMM, and that can be abstracted according to the rules described in section 4 up and until 4.4 of [2]. As our program is not written in C, but in BITCHAKU, and the output language must be recognized by our model checker, Fender, we state the following alterations to the aforementioned rules, there are no procedure calls, pointers, arrays, or any variables of a type other than integer in our input program, and no variables of a type other than boolean in the output. In this section, we define the components, and lastly, the overall structure of our output boolean program. We will use the notation $\sigma[a/b]$ to denote σ with all occurrences of **b** replaced by **a**.

For the predicate abstraction itself, we declare two boolean variables for each predicate $\varphi_i \in E$: B_i and T_i . Both of them hold the truth value of φ_i , albeit at different times. We can guarantee however, that B_i is always the truth value of φ_i . These variables are also declared for the additional predicates defined below. Also, keeping in line with the notation in the paper by Ball et al. [2], we will use the notation $WP(s, \varphi)$ for the weakest liberal precondition of the predicate φ with respect to the statement s , and $\mathcal{F}_V(\varphi)$ for the largest

disjunction of implicative cubes for a predicate φ over a set of boolean variables V .

5.1 Variables and predicates

Our program has a set of variables H , as declared in the original input program, and a set of predicates E over H . Among the newly declared local variables, i.e. all variables that weren't part of the original program, but since have been declared during the buffer size analysis stage, we differentiate between auxiliary variables, which encompass all first pointers and counters of the forms $\mathbf{x_fst_i}$ and $\mathbf{x_cnt_i}$, respectively, and buffer variables, which describe all variables abstracting WMM write buffer data. All buffer variables are assigned to a single global variable at any program point. This follows from the buffer variables' inherent association with a single global variable in the case of TSO, and PSO buffer variables' only being used after allocation to a global variable. With this information, we extend our predicate set by $\mathbf{x_cnt_i} \geq 0$, $\mathbf{x_cnt_i} \leq K$, $\mathbf{x_fst_i} \geq 0$, and $\mathbf{x_fst_i} \leq K$ for each global variable \mathbf{x} and each valid process number i . We will handle these new predicates as we would handle the original ones. Furthermore, we iterate through the program, and for every TSO buffer variable $\mathbf{x_j_i}$ we encounter, we extend our predicate set by $\{\varphi' \mid \varphi' := \varphi[\mathbf{x_j_i}/\mathbf{x}] \wedge \mathbf{x} \in \varphi \in E\}$. We will call our extended predicate set E' . Though we will handle the latter predicates the same way we would handle the others, unless stated otherwise, we will make note of them being *buffer predicates* and copies of certain different predicates $\in E$.

5.2 The choose statement

The `choose(pos, neg);` statement is defined in subsection 4.3. Paraphrased, it states for two boolean variables A and B :

```

if (A  $\equiv$  true)
    choose(A, B)  $\equiv$  true
else if (B  $\equiv$  true)
    choose(A, B)  $\equiv$  false
else
    choose(A, B)  $\equiv$  *
endif;

```

Note that boolean variables in Fender may have the value $*$, i.e. indecisive. We will make use of the `choose` statement during our predicate abstraction of assignments.

5.3 Weakest liberal preconditions

As defined in subsection 4.1 of [2], $WP(s, \varphi)$ is defined as the weakest predicate whose truth before s entails the truth of φ after s terminates (if s terminates), and in the case of s being an assignment statement of the form $\mathbf{x} = \mathbf{e}$, it equals

$\varphi[e/x]$. In this work, we will only compute the weakest liberal preconditions of assignment statements.

5.4 Largest implicative cubes

Paraphrasing the definitions in subsection 4.1 of [2] to more suit the semantics of this work, we define a cube over a boolean variable set V as follows: An implicative cube for a predicate φ is a conjunction of boolean variables $\in V$ (negated or not) representing different predicates $\in E$, which implies φ . We say for a cube C to be the "superset" of another cube D iff D can be expressed as the conjunction of C and a different non-empty cube C' lacking any variable appearing in C . Continuing with our paraphrasing, we define $\mathcal{F}_V(\varphi)$ with the size limit L as the disjunction of all cubes over V implying φ with L or less literals, and $\mathcal{G}_V(\varphi) := \neg\mathcal{F}_V(\neg\varphi)$. For practicality reasons, and without loss of generality, we can redefine the largest disjunction of implicative cubes as lacking any two cubes of which one is the superset of another. We construct such cube disjunctions conservatively, i.e. once we find that a cube c implies a predicate, we will omit any supersets of c from our disjunction, as described in subsection 5.2 of [2].

5.5 Replacement rules

We use the replacement rules defined in section 4 of [2], albeit in a reduced and slightly adapted form, as we need not consider the caveats of the C language.

5.5.1 Local assignment-, store-, and load statements

We reduce all three types of assignment statements to the same form $\mathbf{l} = \mathbf{r}$. First, we examine the left-hand side of the assignment. If it is a buffer variable, it will be regarded as the global variable it represents for all purposes with the exception of determining the left-hand sides of all initialization-, assignment- and reset statements. Next, we build the set of relevant predicates R out of any predicate that contains \mathbf{l} , save for buffer predicates. As long as we encounter predicates containing non-buffer variables, we extend R by all predicates containing the new variables, which will result in the extended set of relevant predicates R' . All the while, we keep ignoring and omitting buffer variables and buffer predicates. Once this closure has been computed, we have all we need to replace the assignment with:

```
begin_atomic;
```

[initialization of temporary variables]

[assignment of boolean variables]

[reset of temporary variables]

```

assume (!( $\mathcal{F}_V(\text{false})$ ));

end_atomic;

```

In this atomic block, we simulate the parallel assignment of the new boolean values to their variables. We define the *initialization of temporary variables* for each predicate $\in R$ with the index j as:

```

load T_j = B_j;

```

Analogously, we define the *reset of temporary variables* as:

```

T_j = 0;

```

Which leaves us with the *assignment of boolean variables*, which is also done for each predicate $\in R$ with the index j , and in accordance to subsection 4.3 of [2] as follows:

```

store B_j = choose( $\mathcal{F}_{R'}(WP(1 = r, \varphi_j))$ ,  $\mathcal{F}_{R'}(WP(1 = r, \neg\varphi_j))$ );

```

If the initial left-hand side 1 was a buffer variable, we scan through every cube disjunction we compute and replace every instance of every boolean variable B_j that represents a predicate featuring 1 with the conjunction of B_j and its buffer predicate counterpart.

5.5.2 Conditionals

We consider conditionals of the form `if (φ) [if block] else [else block] endif;` and replace them as follows, using the extended relevant predicate set R' of all variables featured in φ , computed the same way as described in the previous subsection:

```

if (*) goto else label;

assume ( $\mathcal{G}_{R'}(\varphi)$ );

[if block]

goto end label;

else label: assume ( $\mathcal{G}_{R'}(\neg\varphi)$ );

[else block]

end label: nop;

```


5.5.3 Assertions

For any predicate φ in the assertion clause that doesn't contain a $\text{pc}(i)$ counter, we compute its extended relevant predicate set R' , as described above, and replace those predicates with $\mathcal{F}_{R'}(\varphi)$.

5.5.4 Overall boolean program structure

We build our final program to be verified with our model checker the following way:

```
shared B_0, ..., B_ $|E'|$  - 1;  
local T_0, ..., T_ $|E'|$  - 1;  
  
init  
  store B_0 = *;  
  ...  
  store B_ $|E'|$  - 1 = *;  
  load T_0 = B_0;  
  ...  
  load T_ $|E'|$  - 1 = B_ $|E'|$  - 1;  
  
  [predicate abstractions of initialization stores]  
  
  T_0 = 0;  
  ...  
  T_ $|E'|$  - 1 = 0;  
  
  assume(!( $\mathcal{F}_V(\text{false})$ ));  
  
process 1  
  assume(!( $\mathcal{F}_V(\text{false})$ ));  
  
  [predicate abstraction of the process block]  
  
...  
  
[predicate abstraction of the assertion]
```

6 Discussion

6.1 Complexity

6.1.1 Buffer size analysis

Computing the buffer size increases caused by each program point has a time complexity linear to the number of program points. Since computing the actual

buffer sizes needed at each program point requires a comparison of all buffer sizes per variable per program point. Therefore, the overall time complexity for the buffer size analysis in itself is $\mathcal{O}(mn)$, with m being the number of program points in P , and n being the number of global variables in P .

As for the replacements carried out in the buffer size analysis stage, we consider the most complex WMM operation that will be replaced at this stage: the flush statement. Flush statements generate a separate flush code block for every global variable, and the flush code block for any variable x contains a maximum of $\mathcal{O}(K)$ number of nested if-else-clauses.

Therefore, the buffer size analysis stage has a time complexity of $\mathcal{O}(\max(mK, mn))$

6.1.2 Predicate abstraction

6.2 Evaluation

We have implemented the procedures as described so far in a tool we call SRO-TOGAP. Then, we have run numerous tests on the following concurrent algorithms: the alternating bit protocol, Lamport’s Bakery Algorithm, Dekker’s Algorithm, Peterson’s Algorithm, a lock-free queue, Szymański’s Mutual Exclusion Algorithm, and the ticket locking algorithm. All experiments were conducted on an Intel(R) Core(TM) i7-3770 3.40GHz with 8GB RAM. The key question was whether

References

- [1] A. Dan, Y. Meshman, M. Vechev, and E. Yahav. Predicate Abstraction for Relaxed Memory Models. In Static Analysis Symposium ’13 (2013)
- [2] T. Ball, R. Majumdar, T. Millstein, and S. K. Rajamani. Automatic Predicate Abstraction of C Programs. In PLDI ’01 (2001)
- [3] Fender. Model checking, abstract interpretation, and dynamic analysis. ETH Zurich and Technion. <http://www.practicalsynthesis.org/fender/>